

Foundations

- Searching Basics
 - Searching is most common DB operation
 - SELECT - most versatile
- Data Storage Structures
 - Arrays: Fast random access, slow insertions
 - Linked Lists: Fast insertions, slow random access
- Database Search Challenges
 - Data stored by ID is fast to search.
 - Searching by other attributes requires linear scan (inefficient)
- Optimizing Search & Inserts
 - Sorted arrays allow binary search but slow inserts
 - Linked lists allow fast inserts but slow searches
 - Binary Search Trees (BSTs) balance fast inserts and searches

BST URochester

- A binary tree where each node has a key, and for any node x, all keys in its left subtree are less than x's key, and all keys in its right subtree are greater
- Traversal Methods:
 - Inorder Traversal (Visits the left subtree, then the current node, and finally the right subtree. In a BST, this traversal retrieves keys in ascending order), Preorder Traversal (before sub), Postorder Traversal (after sub)
- Search → Compare with root, move left if smaller, right if larger
- Find Min/Max →
 - Min: Go left until no more left child
 - Max: Go right until no more right child
- Insertion → Find correct position, insert while maintaining order
- Successor →
 - If right child exists → Leftmost node in right subtree
 - Otherwise → First ancestor where node is in the left subtree
- Predecessor →
 - If left child exists → Rightmost node in left subtree
 - Otherwise → First ancestor where node is in the right subtree

AVL UCI Notes

- AVL: self balancing
- Insertion: insert like BST, then rebalance if unbalanced
- Deletion: remove node, rebalance all the way up
- Search: same as BST but always $O(\log n)$
- Single Rotation: Fixes left-left or right-right imbalance
- Double Rotation: Fixes left-right or right-left imbalance

AVL Rotations:

- Right Rotation: Fixes left-heavy imbalance
- Left Rotation: Fixes right-heavy imbalance
- Left-Right Rotation: First left rotate child, then right rotate parent
- Right-Left Rotation: First right rotate child, then left rotate parent

B Trees:

- Self-balancing tree: Keeps data sorted for fast operation
- Search, insert, delete in $O(\log n)$ time
- All leaf nodes are at the same level
- Each node holds $m/2$ to m children (except root)
- Start from root → insert in appropriate leaf (If full, split the node and push the middle key up) → deletion (remove the key, if a node underflows, borrow from siblings or merge nodes) → balancing (Nodes are split or merged to keep the tree balanced)

B+ Trees:

- Insertion: Start at the root and follow keys to the correct leaf node.
 - If leaf is full, split it, move half the values to a new node, and copy the smallest key to the parent
- Splits:
 - Leaf split → Copy smallest key to parent
 - Internal split → Move middle key up
 - Root split → Adds a new level to the tree

Mod 3 Moving Beyond Relational

- Relational Databases → Strong ACID compliance, great for structured data, but scaling is hard
- Challenge: Schema changes, expensive joins, data limitations

Mod 5 NoSQL

- NoSQL: Scales horizontally, handles unstructured data, uses BASE instead of ACID
- CAP Theorem: Can only pick two: Consistency, Availability, or Partition Tolerance
- Key-Value Stores: Fast ($O(1)$) lookups, simple key = value structure, no joins
- Use Cases:
 - Session storage
 - Caching
 - Feature stores
 - Shopping carts
- Redis: In-memory KV store, supports multiple data types, ultra-fast (>100K ops/sec)

Mod 6 Redis Python

- `redis.Redis(host, port, db, decode_responses=True)`
- Strings → `set()`, `get()`, `incr()`, `mset()`.
- Lists → `rpush()`, `lpush()`, `lrange()`, `lpop()`.
- Hashes → `hset()`, `hgetall()`, `hdel()`

Mod 7 Doc DBs

- MongoDB
 - Stores data as JSON/BSON
 - indexing, replication, load balancing
 - Similar to SQL

Mod 8 PyMongo

- Insert → `insert_one()`, `insert_many()`.
- Find → `find()`, `find_one()`.
- Update → `update_one()`, `update_many()`.

- Delete → `delete_one()`, `delete_many()`